

# CS533 Project Progress: Map Reduce

*Abdul Dakkak, Carl Pearson, Liwen Chang*

## Project Summary

Our project proposal stated that we will be implementing a compiler for a functional language that would generate MPI code that would run on a cluster of machines and across threads in a shared memory system. We have narrowed down our objective of generating code that would run on the GPU. The interesting features in the compiler and language, as stated in the proposal, is that parallelization will be done via Map and Reduce operations. The compiler will maintain the IR for these operators and will perform optimizations on them to generate performant code.

It is worth mentioning that our aim is not to write an optimized compiler (the compiler can be very slow), our aim is to have a compiler that can generate optimized code. To that end, we have written the compiler in Dart (a Javascript inspired language) that currently generates sequential Javascript code from our language. The Javascript generation is meant primarily for debugging purposes and the next steps would be to generate threaded CUDA code and perform some compiler optimizations on the IR. The next few paragraphs revise and summarize the project as well as detail what has been implemented and what will be implemented. At the end we give a schedule for the project.

## Language Details

Our statically typed language is inspired by array and data flow programming languages such as Fortran[1], APL[2], and LINQ[3] where one expresses computation based on operations on vectors and arrays. In our language, one can define two vectors of size 100 by

```
n :: Integer = 1000;
as :: []Real = rand.Real(n);
bs :: []Real = rand.Real(n);
```

We can then add the two vectors using an overloaded plus operator

```
res :: []Real = as + bs;
```

To give you a taste of the language, in this program we approximate pi using Monte Carlo integration

```
def f(a :: Real, b :: Real) :: Bool {
  return a*a + b*b < 1;
}
res :: Integer = zip(as, bs).count(f) / n;
```

this would be translated into the map/reduce operations of

```
t1 = map(f, zip(as, bs));
count = 0;
reduce((x) => count += x, t1);
res = count / n;
```

The compiler would lower the above into an IR representation, maintaining the map/reduce operations

```
Instruction(zab, zip(as, bs))
Instruction(t1, Map(f, zab))
Instruction(count, 0)
Function(g, (x) => count += x)
Instruction(r2, Reduce(g, t1))
Instruction(res, divide(count, n))
```

The compiler is able to analyze the above code, inline the map operation into the reduce function and privatize the count variable. This results in efficient code that can be parallized and does not produce unnecessary temporary arrays.

We plan on expressing all commonly used analytics operations such as sort, mean, max, min, histogram, variance, etc. . . in this framework.

## Compiler Pass

The most important factor in distributed computing is how to manage memory transfer. If a node computes a chunk of data and it is used in subsequent instructions, then it should reuse the output rather than send and request the data again. There are two approaches to facilitate this. The first is a runtime approach: Hadoop, for example, dispatches tasks to maximize reuse of local data. This done via the Hadoop scheduler which has a mapping between nodes and data state.

The second is a compiler transformation. This is mainly done via loop fusion. If for example, one writes a program `map(f, map(g, lst))` then a compiler pass can transform this into `map(f.g, lst)`. A simple peephole optimizer can scan for this instruction pattern and perform this transofmration. A generalization of this technique for other list primitives is found in the the Haskell vector library. Using a concept called Stream Fusion[4], Haskell fuses most function loops to remove unnecessary temporaries and list traversals. In this project, we will adopt some aspects of how Haskell performs this transformation when they are applicable in a CUDA programming model.

## Implementation Details

To facilitate rapid prototyping in this project, we chose the Dart programming language. The Dart language developed by google is a modern interperetation of Javascript — a cross between C++/Java and Javascript. It adds classes, types, and polymorphic instances (via a templating mechanism) and is able to compile down to Javascript or can be run in the Dart VM. The language also has good documentation, extensive standard library, and an active library development community (the parser generator, for example, is a library we are using).

The way we structured our compiler, it is backend agnostic. Currently, we generate sequential Javascript for debugging purposes, but have a prototype CUDA backend that generates naive code. In the next few weeks, we plan on refining our CUDA implementation to hide memory copy latency and optimize for the correct launch paramaters for the kernel.

We will use parboil as our benchmark suite, picking 4-5 benchmarks that map nicely to our language. We will then measure the performance obtained from our compiler versus hand optimized GPU code. We will also compare the programming difficulty in both programming interfaces.

## Progress Summary

We have

- Parser

- Instruction lowerer
- Backend to generate CUDA code
- Backend to generate Javascript code (this is used for debugging)

## Project Schedule

We have developed the infrastructure to allow us to start working on the interesting parts of the project. In the rest of the semester, we will be concentrating on developing compiler passes that would allow us to generate efficient backend code. The following table is our projected timeline for the rest of the semester.

Week	Task
3/17	Finish naive CUDA code generator.
3/24	Add compiler pass to perform closure conversion (for lambda functions) and calculate the def-use chain of variables.
3/31	Generate optimized map kernels (this requires finding tuning parameters for architectures, but NVIDIA provides tools to programatically determine those parameters).
4/07	Generate optimized reduce kernels (this, again, requires some tuning, but a group member has done extensive research on reduce operations on GPUs).
4/14	Add compiler pass to perform function fusion.
4/21	Experiment with other compiler passes, such as loop unrolling, that would increase the compute work done by each thread.
4/28	Final benchmarking and project writeup.
5/05	Complete project presentation.

Projected timeline for the project along with the associated tasks.

## References

- [1] J.C. Adams, W.S. Brainerd, and C.H. Coldberg, *Programmer's guide to Fortran 90*, Intertext Publications, 1990.
- [2] R.P. Polivka and S. Pakin, *APL: The language and its usage*, Prentice Hall Professional Technical Reference, 1975.
- [3] E. Meijer, B. Beckman, and G. Bierman, "Linq: reconciling object, relations and xml in the .net framework," *Proceedings of the 2006 ACM SIGMOD international conference on Management of data*, ACM, 2006, pp. 706–706.

[4] D. Coutts, R. Leshchinskiy, and D. Stewart, "Stream fusion: From lists to streams to nothing at all," *ACM SIGPLAN Notices*, ACM, 2007, pp. 315–326.